

```
In [140]: import pandas as pd
```

```
In [141]: import numpy as np
```

```
In [142]: import seaborn as sns
```

```
In [143]: import matplotlib.pyplot as plt
```

```
In [144]: %matplotlib inline
```

```
In [145]: from sklearn.linear_model import LogisticRegression
```

```
In [146]: from sklearn.externals import joblib
```

```
In [147]: diabetesDF = pd.read_csv('DESKTOP/diabetes.csv')
```

```
In [148]: print(diabetesDF.head())
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI \
0	6	148	72	35	0	33.6
1	1	85	66	29	0	26.6
2	8	183	64	0	0	23.3
3	1	89	66	23	94	28.1
4	0	137	40	35	168	43.1

	DiabetesPedigreeFunction	Age	Outcome
0	0.627	50	1
1	0.351	31	0
2	0.672	32	1
3	0.167	21	0
4	2.288	33	1

In [149]: `diabetesDF.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
Pregnancies      768 non-null int64
Glucose          768 non-null int64
BloodPressure    768 non-null int64
SkinThickness    768 non-null int64
Insulin          768 non-null int64
BMI              768 non-null float64
DiabetesPedigreeFunction  768 non-null float64
Age              768 non-null int64
Outcome          768 non-null int64
dtypes: float64(2), int64(7)
memory usage: 54.1 KB
```

In [150]: `corr = diabetesDF.corr()`

In [151]: `print(corr)`


	Pregnancies	Glucose	BloodPressure	SkinThi
ckness \				
Pregnancies	1.000000	0.129459	0.141282	-0.
081672				
Glucose	0.129459	1.000000	0.152590	0.
057328				
BloodPressure	0.141282	0.152590	1.000000	0.
207371				
SkinThickness	-0.081672	0.057328	0.207371	1.
000000				
Insulin	-0.073535	0.331357	0.088933	0.
436783				
BMI	0.017683	0.221071	0.281805	0.
392573				
DiabetesPedigreeFunction	-0.033523	0.137337	0.041265	0.
183928				
Age	0.544341	0.263514	0.239528	-0.
113970				

Outcome	0.221898	0.466581	0.065068	0.
074752				
	Insulin	BMI	DiabetesPedigreeFunction	
\				
Pregnancies	-0.073535	0.017683		-0.033523
Glucose	0.331357	0.221071		0.137337
BloodPressure	0.088933	0.281805		0.041265
SkinThickness	0.436783	0.392573		0.183928
Insulin	1.000000	0.197859		0.185071
BMI	0.197859	1.000000		0.140647
DiabetesPedigreeFunction	0.185071	0.140647		1.000000
Age	-0.042163	0.036242		0.033561
Outcome	0.130548	0.292695		0.173844

	Age	Outcome
Pregnancies	0.544341	0.221898
Glucose	0.263514	0.466581
BloodPressure	0.239528	0.065068
SkinThickness	-0.113970	0.074752
Insulin	-0.042163	0.130548
BMI	0.036242	0.292695
DiabetesPedigreeFunction	0.033561	0.173844
Age	1.000000	0.238356
Outcome	0.238356	1.000000

```
In [152]: sns.heatmap(corr,xticklabels=corr.columns, yticklabels=corr.columns)
```

```
Out[152]: <matplotlib.axes._subplots.AxesSubplot at 0x18305bc7780>
```



```
In [153]: dfTrain = diabetesDF[:650]
```

```
In [154]: dfTest = diabetesDF[650:750]
```

```
In [155]: dfCheck = diabetesDF[750:]
```

```
In [156]: trainLabel = np.asarray(dfTrain['Outcome'])
```

```
In [157]: trainData = np.asarray(dfTrain.drop('Outcome',1))
```

```
In [158]: testLabel = np.asarray(dfTest['Outcome'])
```

```
In [159]: testData = np.asarray(dfTest.drop('Outcome',1))
```

```
In [160]: means = np.mean(trainData, axis=0)
```

```
In [161]: stds = np.std(trainData, axis=0)
```

```
In [162]: trainData = (trainData - means)/stds
```

```
In [163]: testData = (testData - means)/stds
```

```
In [164]: diabetesCheck = LogisticRegression()
```

```
In [165]: diabetesCheck.fit(trainData, trainLabel)
```

```
C:\Users\admin\Anaconda3\lib\site-packages\sklearn\linear_model\logistic.py:433: FutureWarning: Default solver will be changed to 'lbfgs' in 0.22. Specify a solver to silence this warning.  
FutureWarning)
```

```
Out[165]: LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,  
                             intercept_scaling=1, max_iter=100, multi_class='warn',  
                             n_jobs=None, penalty='l2', random_state=None, solver='warn',  
                             tol=0.0001, verbose=0, warm_start=False)
```

```
In [166]: accuracy = diabetesCheck.score(testData, testLabel)
```

```
In [167]: print("accuracy = ", accuracy * 100, "%")
```

```
accuracy = 78.0 %
```

```
In [168]: coeff = list(diabetesCheck.coef_[0])
```

```
In [169]: features = pd.DataFrame()
```


```
In [170]: features['importance'] = coeff
```

```
In [171]: features.sort_values(by=['importance'], ascending=True, inplace=True)
```

```
In [172]: features['positive'] = features['importance'] > 0
```


```
In [173]: features.importance.plot(kind='barh', figsize=(11, 6), color = features.  
positive.map({True: 'blue', False: 'red'}))
```

```
Out[173]: <matplotlib.axes._subplots.AxesSubplot at 0x18305d1d198>
```



```
In [174]: plt.xlabel('Importance')
```

```
Out[174]: Text(0.5, 0, 'Importance')
```



```
In [175]: joblib.dump([diabetesCheck, means, stds], 'diabetesModel.pkl')
```

```
Out[175]: ['diabetesModel.pkl']
```

```
In [176]: diabetesLoadedModel, means, stds = joblib.load('diabetesModel.pkl')
```

```
In [177]: accuracyModel = diabetesLoadedModel.score(testData, testLabel)
```

```
In [178]: print("accuracy = ", accuracyModel * 100, "%")
```

```
accuracy = 78.0 %
```

```
In [179]: print(dfCheck.head())
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI
750	4	136	70	0	0	31.2
751	1	121	78	39	74	39.0
752	3	108	62	24	0	26.0

753	0	181	88	44	510	43.3
754	8	154	78	32	0	32.4

	DiabetesPedigreeFunction	Age	Outcome
750	1.182	22	1
751	0.261	28	0
752	0.223	25	0
753	0.222	26	1
754	0.443	45	1

```
In [180]: sampleData = dfCheck[:1]
```

```
In [181]: sampleDataFeatures = np.asarray(sampleData.drop('Outcome',1))
```

```
In [182]: sampleDataFeatures = (sampleDataFeatures - means)/stds
```

```
In [183]: predictionProbability = diabetesLoadedModel.predict_proba(sampleDataFeatures)
```

```
In [184]: prediction = diabetesLoadedModel.predict(sampleDataFeatures)
```

```
In [185]: print('Probability:', predictionProbability)
```

```
Probability: [[0.4385153 0.5614847]]
```

```
In [186]: print('prediction:', prediction)
```

```
prediction: [1]
```